

A new Ad-Hoc parallel file system for HPC environments based on the Expand parallel file system

Félix García-Carballeira

felix.garcia@uc3m.es

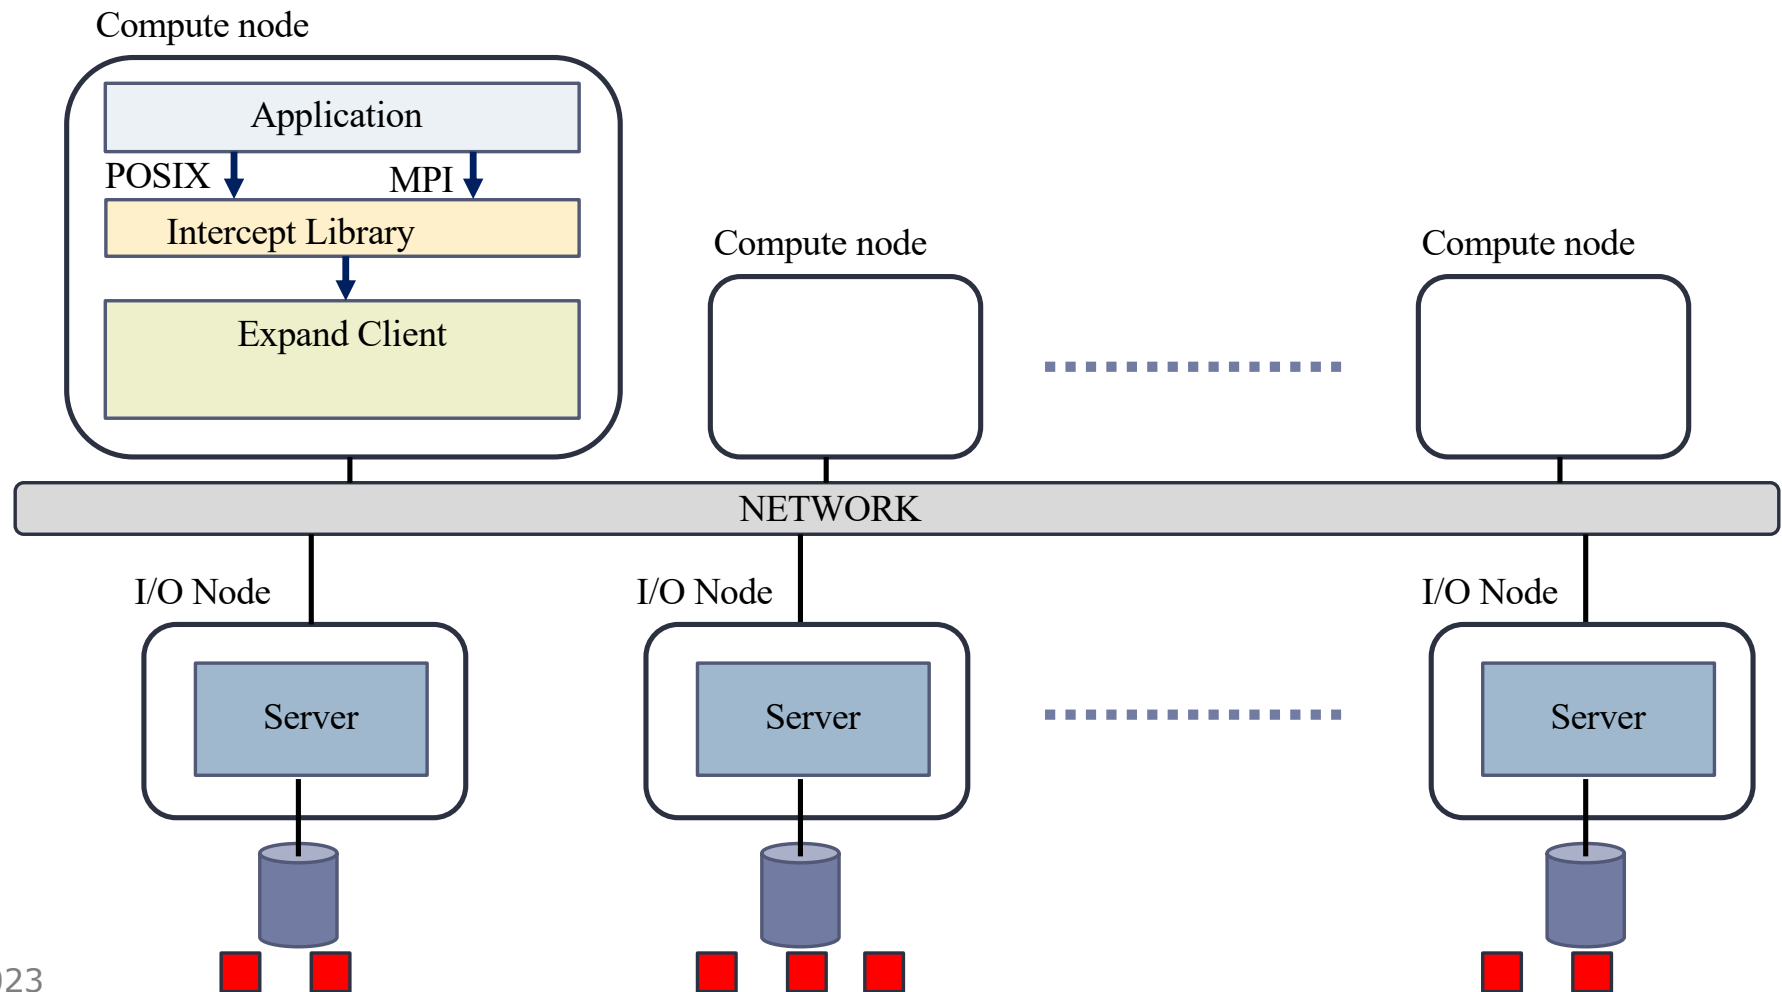


uc3m | Universidad **Carlos III** de Madrid

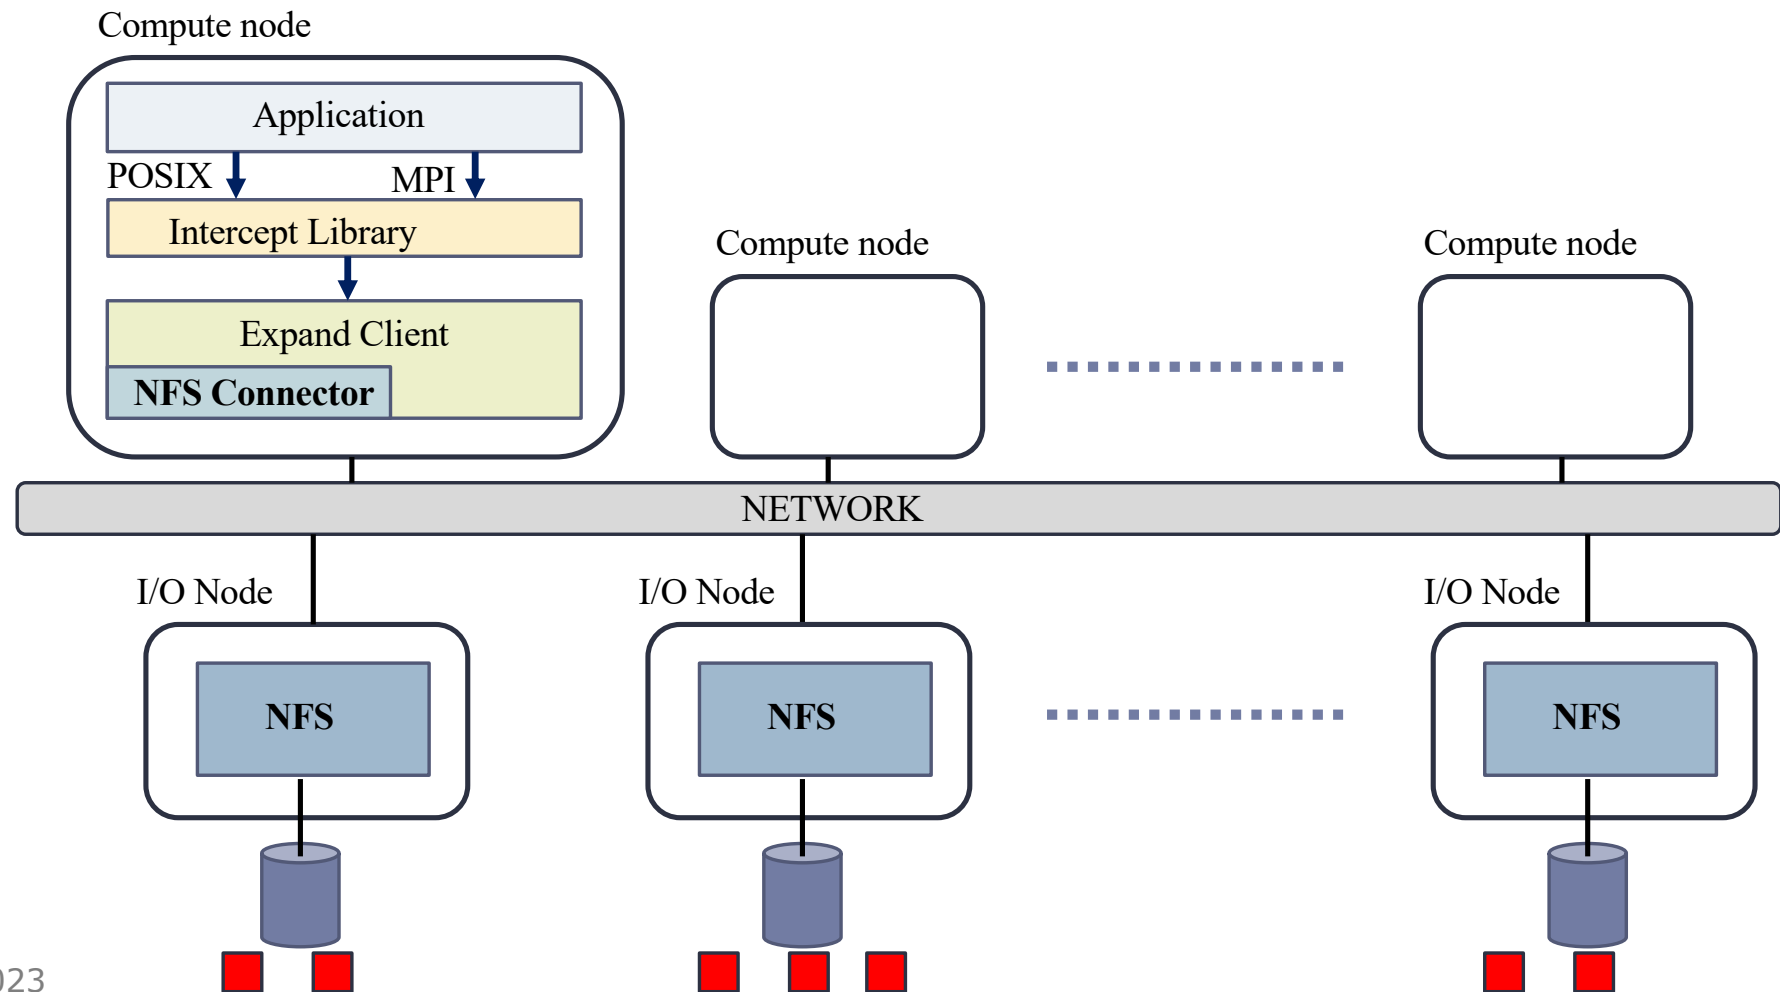
Agenda

- ▶ Introduction to Expand Parallel file system
- ▶ Expand as Ad-hoc parallel file system for large clusters and supercomputers
 - ▶ Motivation
 - ▶ Architecture
 - ▶ Data distribution
 - ▶ Metadata Management
- ▶ Performance evaluation
- ▶ Conclusions and future work

Introduction to Expand

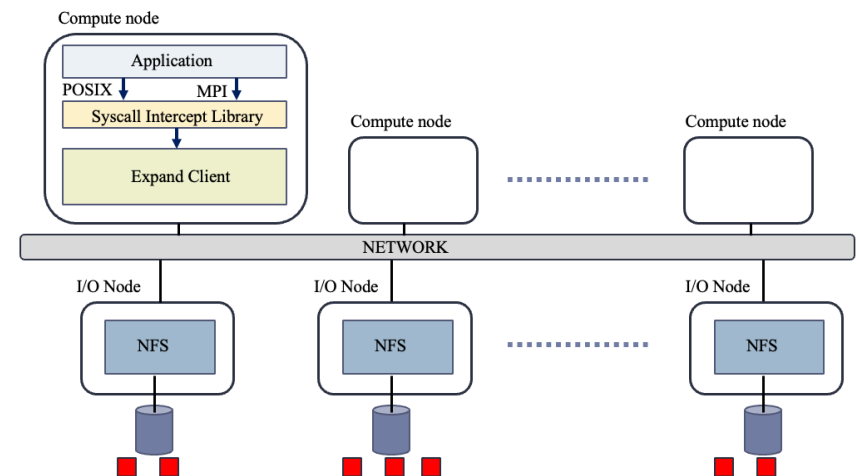


Expand: distributed systems and small clusters

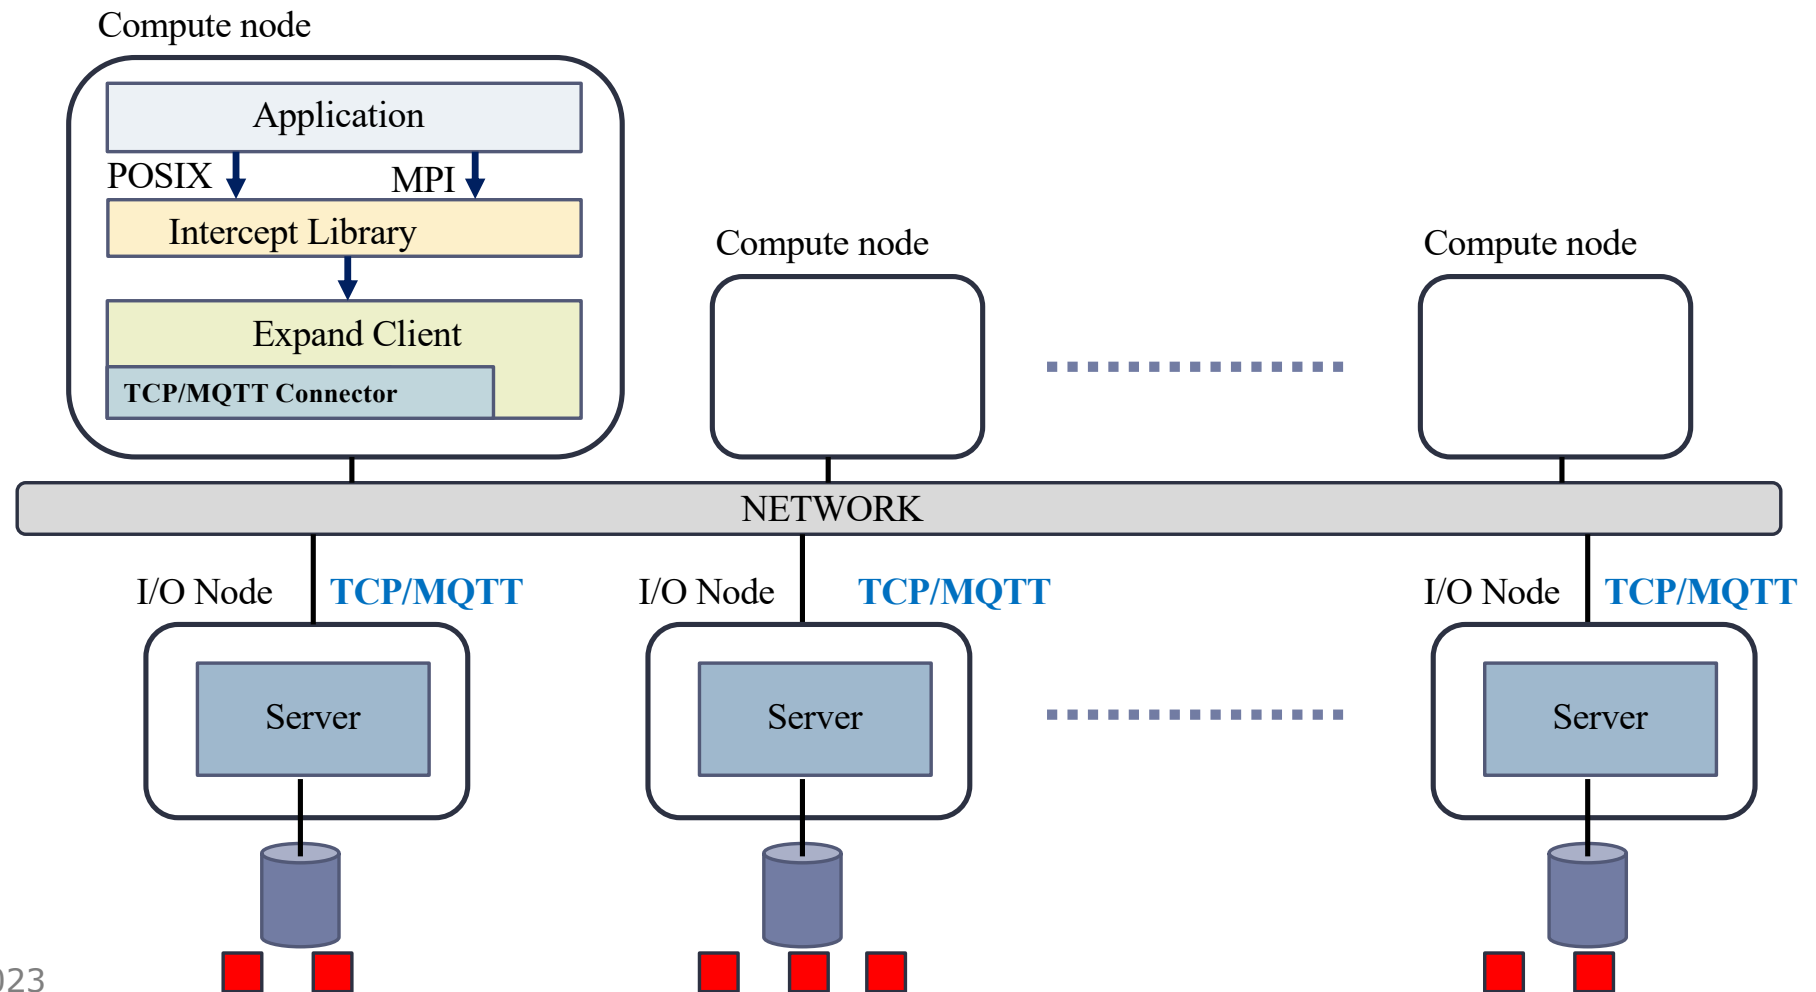


Expand: distributed systems and small clusters

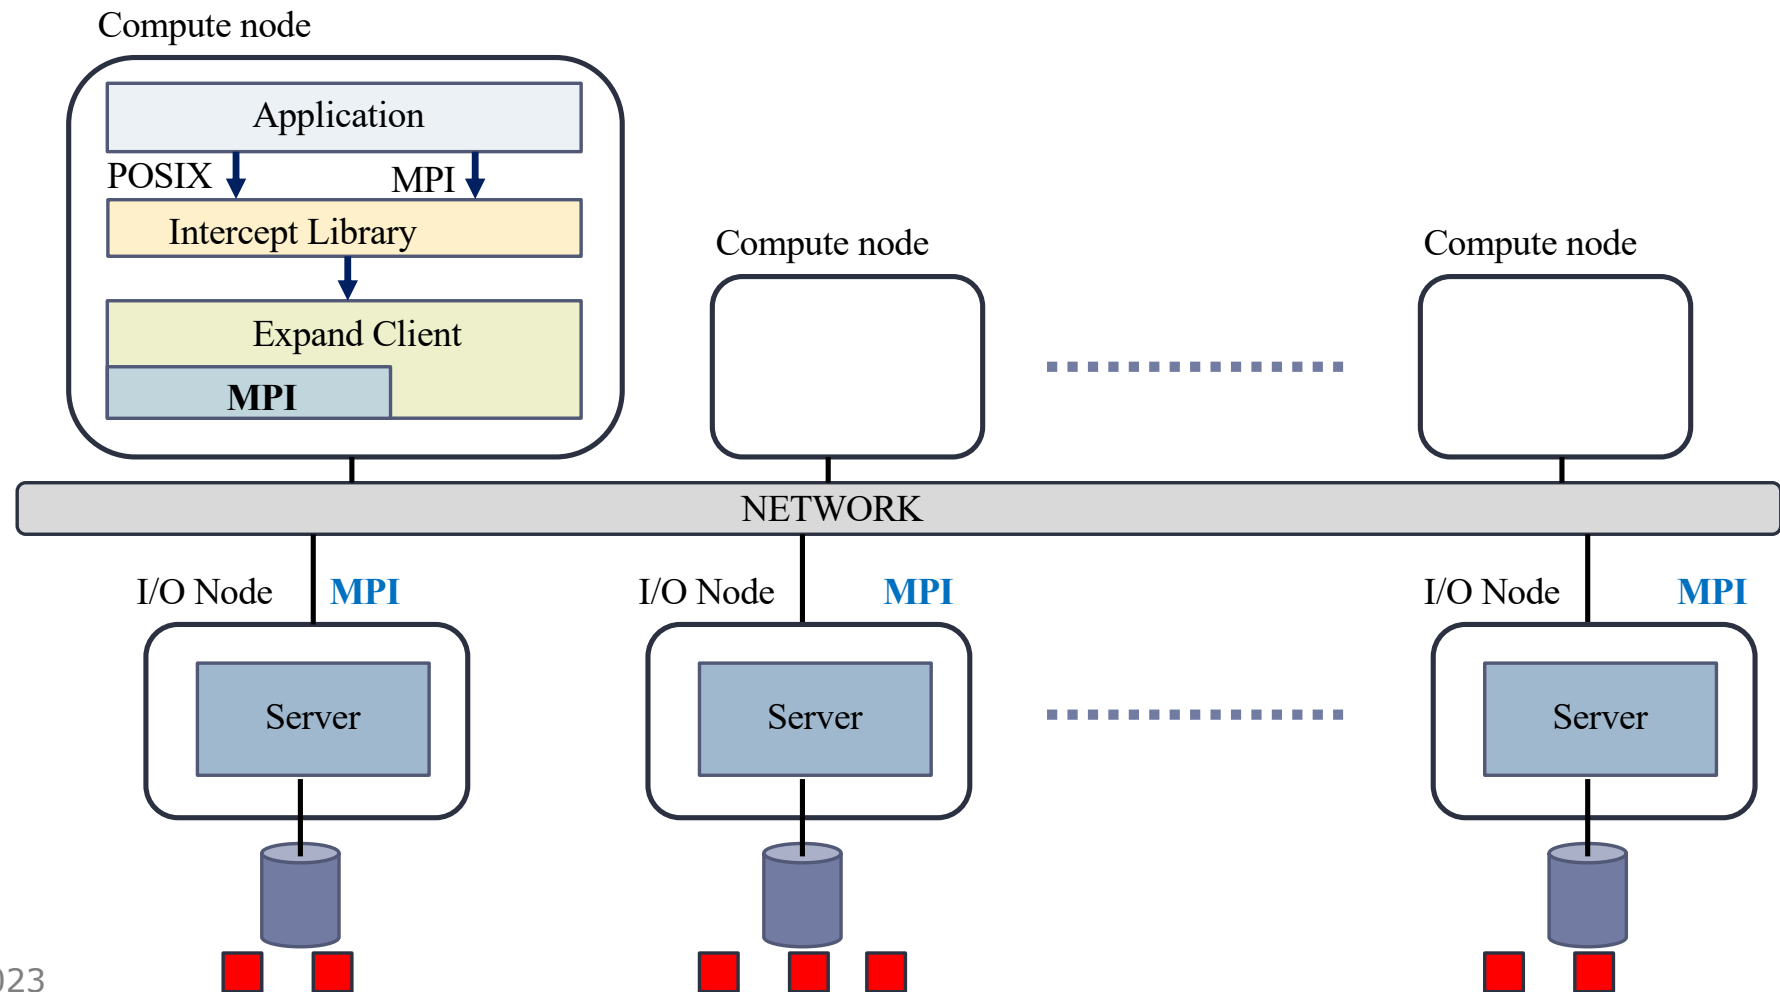
- ▶ Advantages of using NFS:
 - ▶ Standard server
 - ▶ No changes to NFS server are required
 - ▶ Independent of the operating system used in the clients
 - ▶ Allows the using of servers with different architectures and operating systems
 - ▶ The parallel file system construction is greatly simplified
 - ▶ NFS is very familiar to users: easy configuration



Expand: IoT environments

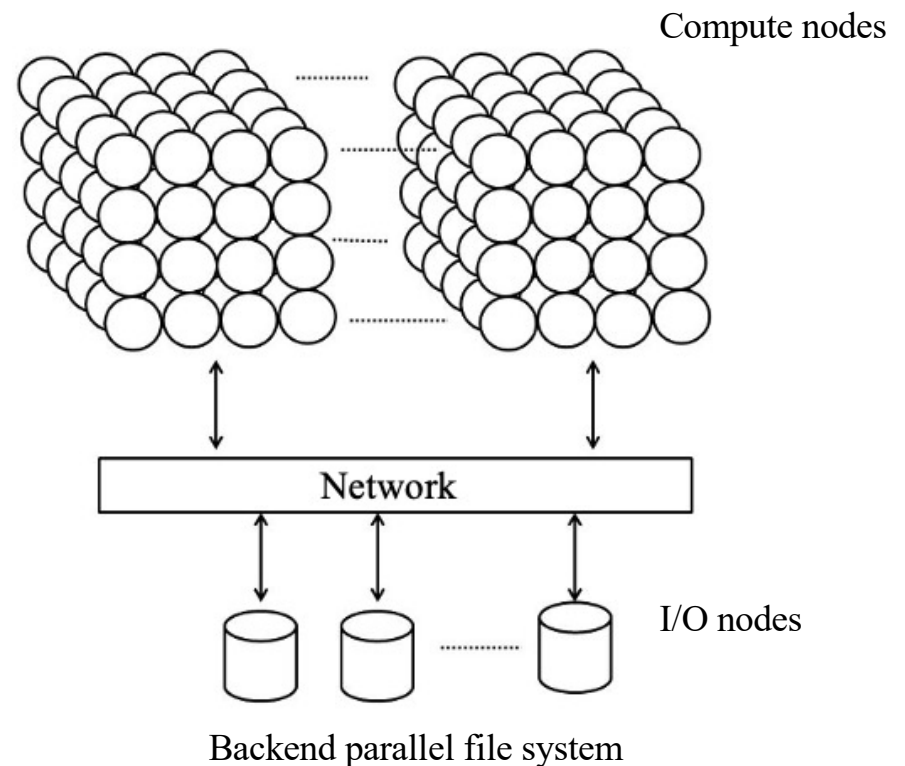


Expand: ad-hoc parallel file system for large clusters




Motivation for developing an Ad-hoc file system

- ▶ Typical supercomputer architecture:
 - ▶ The number of compute nodes is much larger than the number of I/O nodes:
 - ▶ Possible bottleneck



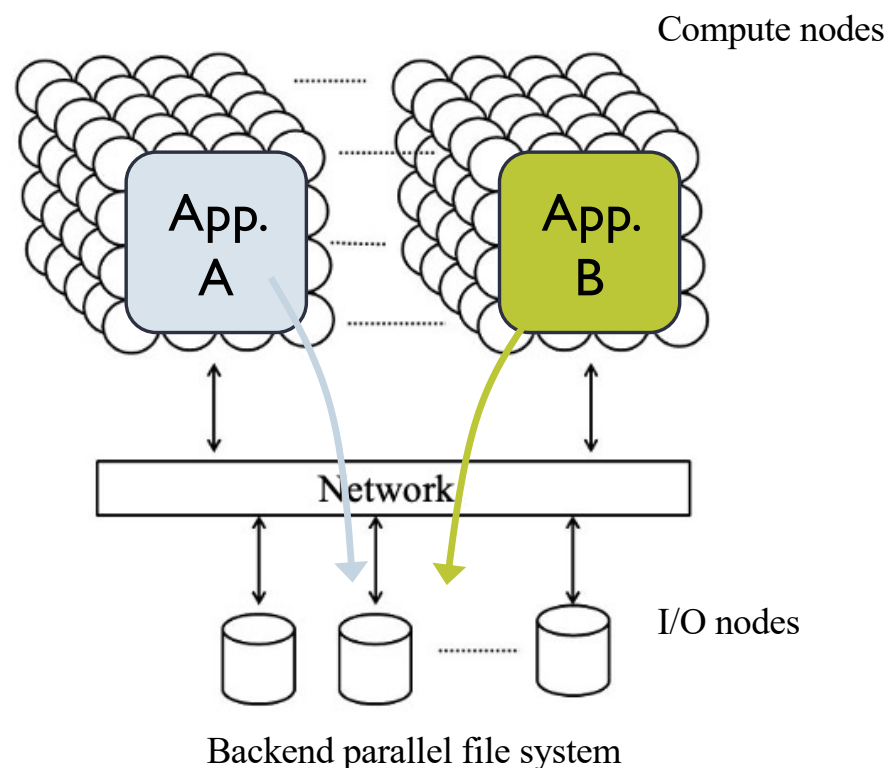
Motivation for developing an Ad-hoc file system

- ▶ Typical supercomputer architecture:
 - ▶ The number of compute nodes is much larger than the number of I/O nodes:
 - ▶ Possible bottleneck
 - ▶ Data away from applications:
 - ▶ Use of network
 - ▶ Reduces data access performance

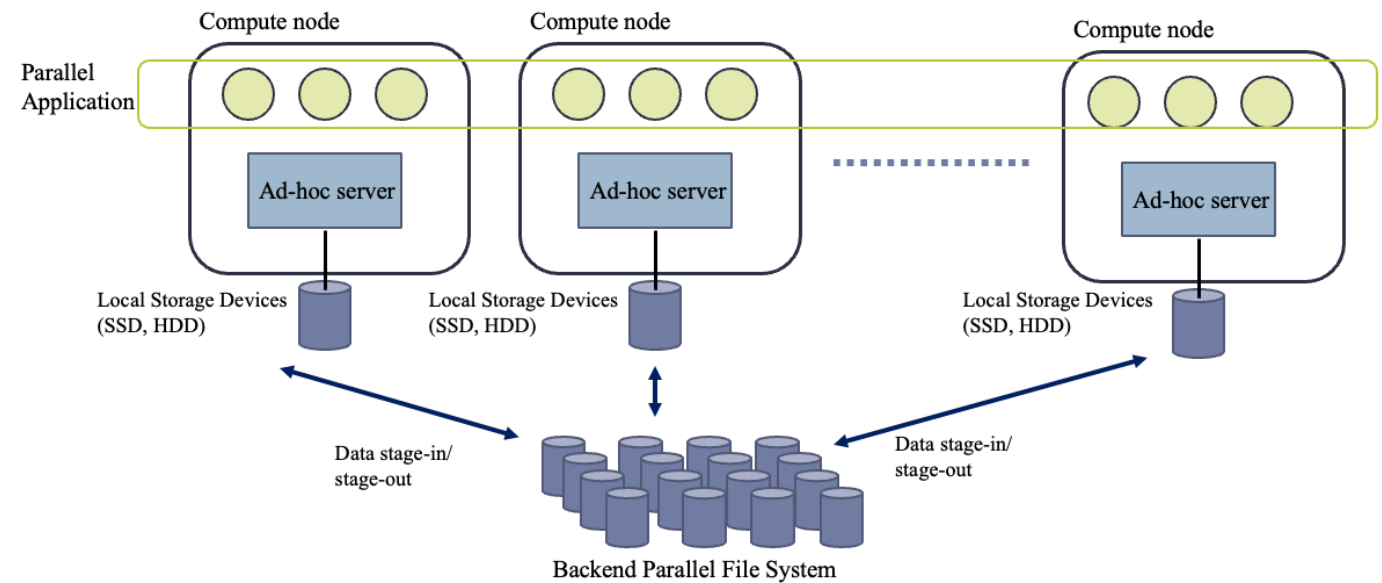


Motivation for developing an Ad-hoc file system

- ▶ Typical supercomputer architecture:
 - ▶ The number of compute nodes is much larger than the number of I/O nodes:
 - ▶ Possible bottleneck
 - ▶ Data away from applications:
 - ▶ Use of network
 - ▶ Reduces data access performance
 - ▶ Different applications running at the same time:
 - ▶ File system conflicts and interferences

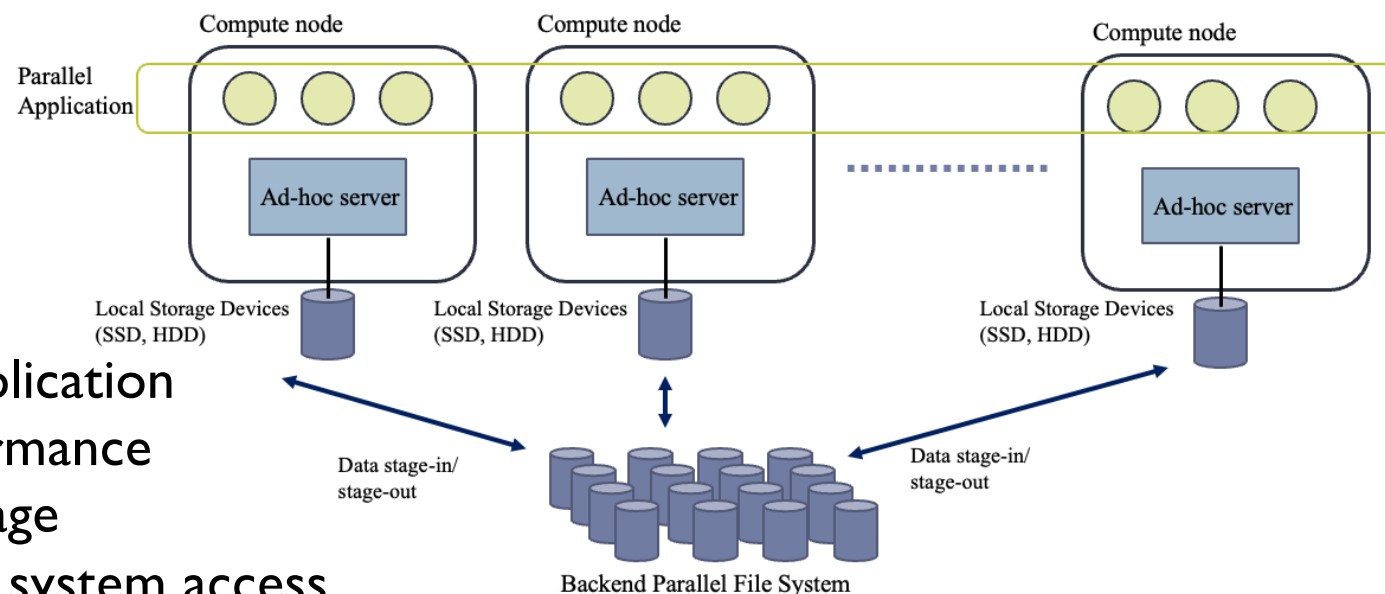


Ac-hoc file systems



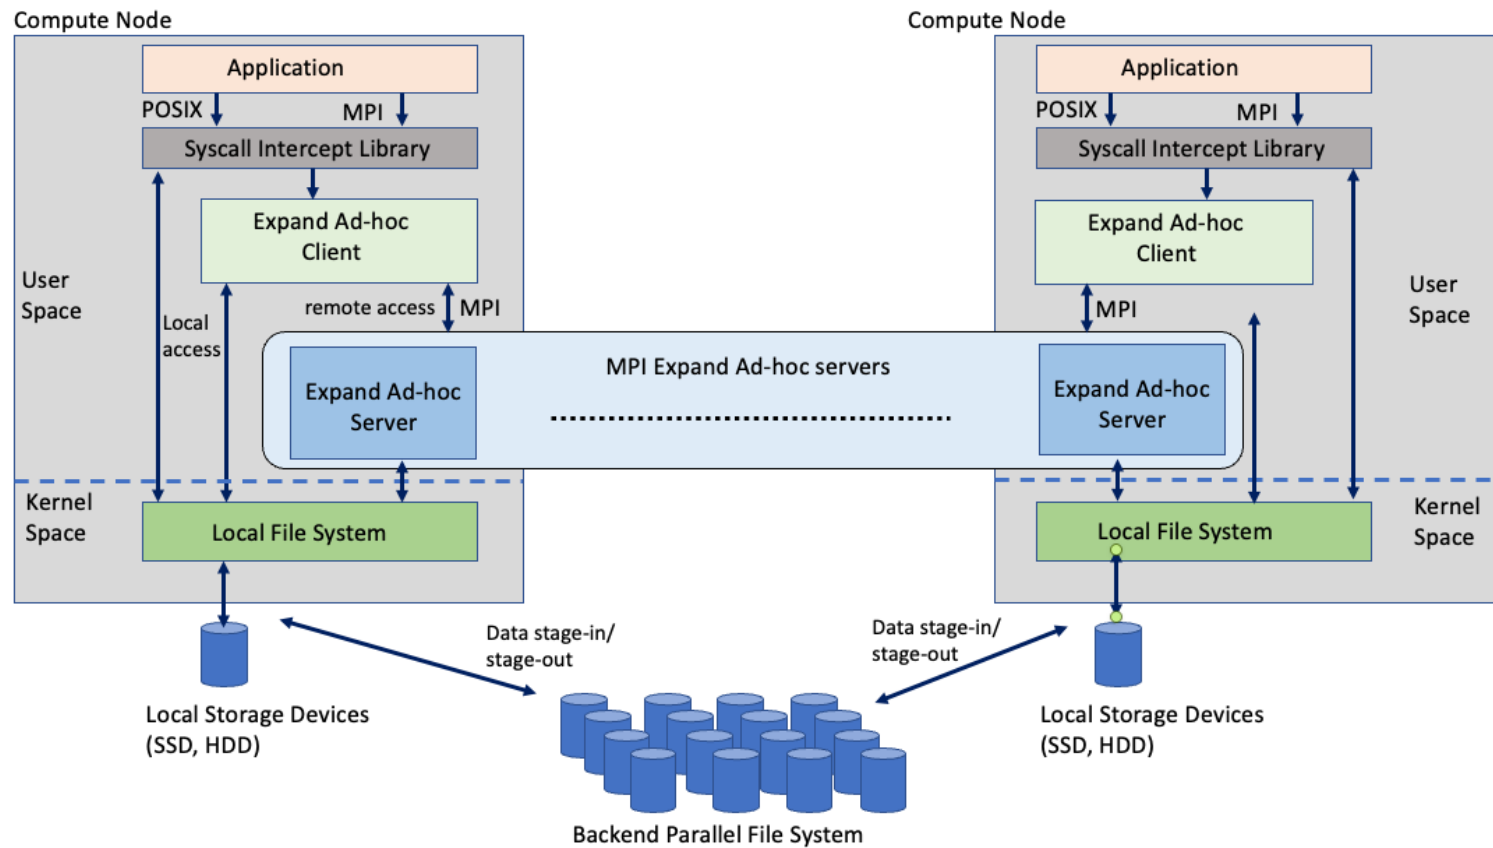
- ▶ Temporary storage system for single applications or workflows of applications that use the available storage of compute nodes (HDD, SSD or SHM)
- ▶ The servers are deployed on the application's compute nodes

Ac-hoc file systems: **advantages**



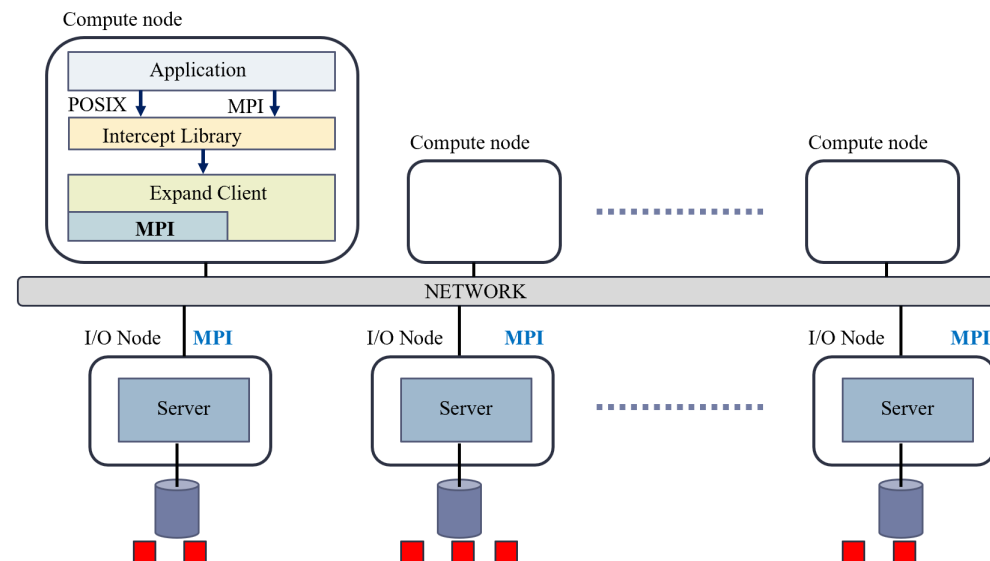
- ▶ Data close to the application
 - ▶ Increase the performance
- ▶ Reduces network usage
- ▶ Reduces backend file system access
 - ▶ Reduces bottlenecks
- ▶ Temporary data (checkpoints, intermediate data) does not need to be stored in the backend file system
- ▶ In application workflows data remains for the different workflow processes

Architecture of the Expand Ad-Hoc



Architecture of the Expand Ad-Hoc

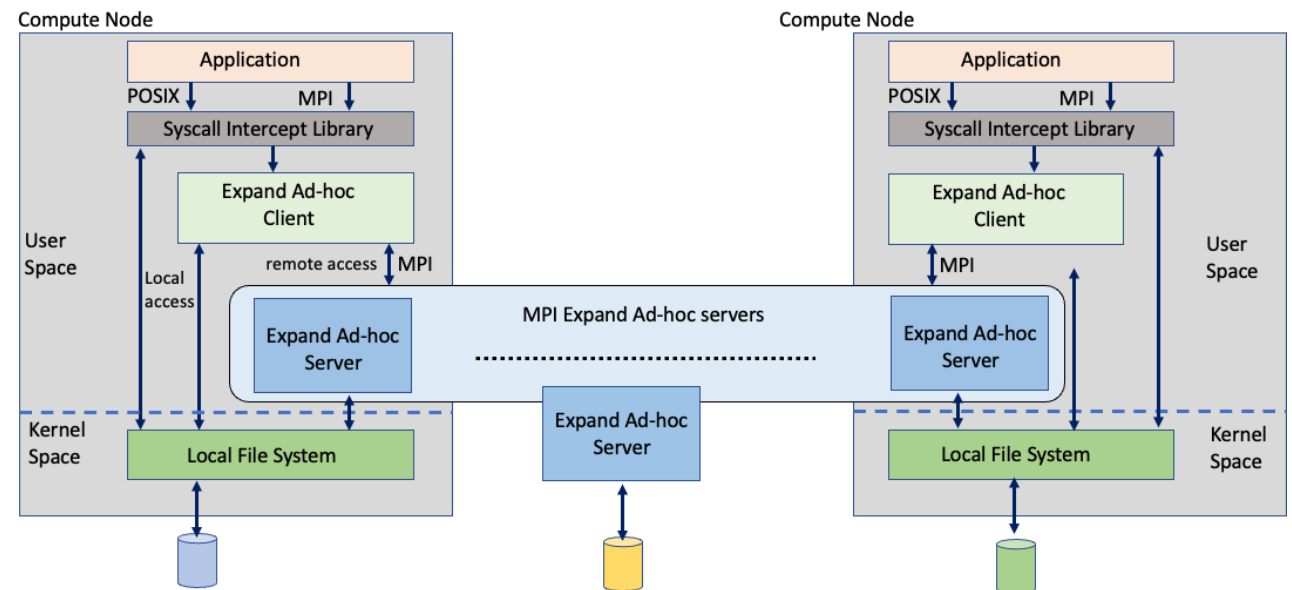
- ▶ Advantages of using MPI:
 - ▶ Standard interface and portable to different platforms
 - ▶ Offers good performance in clusters



Data distribution

- ▶ File in Expand:
 - ▶ Metadata subfile
 - ▶ Several data subfiles
- ▶ Data distributed on different servers

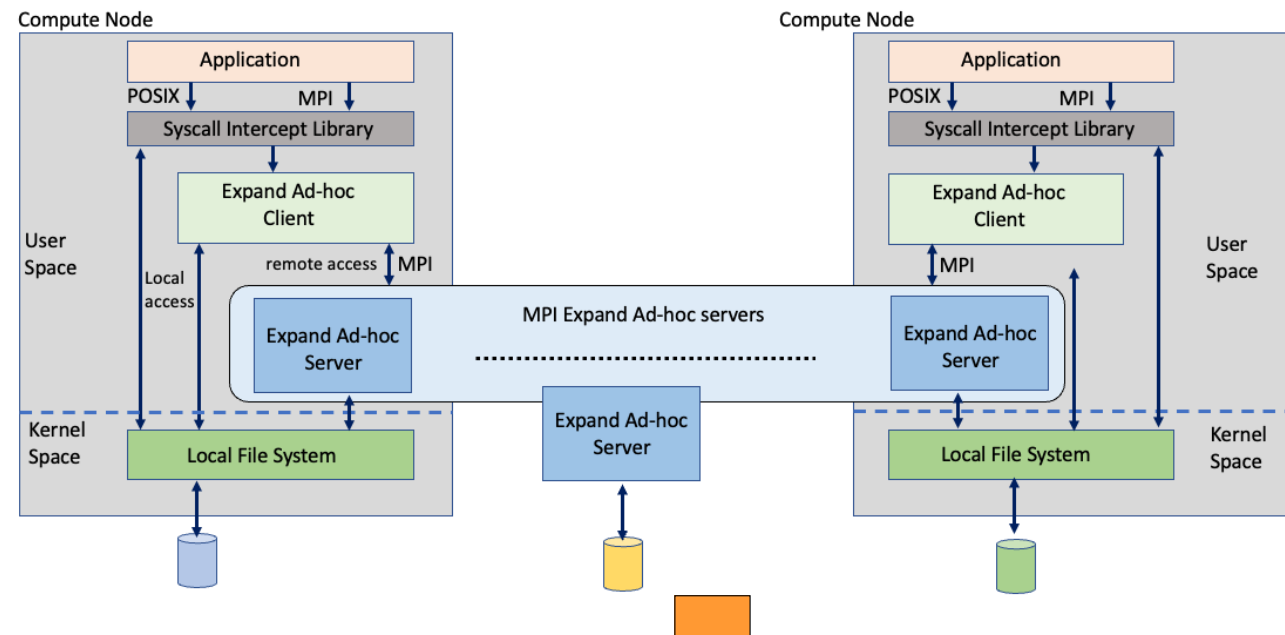
File in Expand



Data distribution

- ▶ File in Expand:
 - ▶ Metadata subfile
 - ▶ Several data subfiles
- ▶ Data distributed on different servers

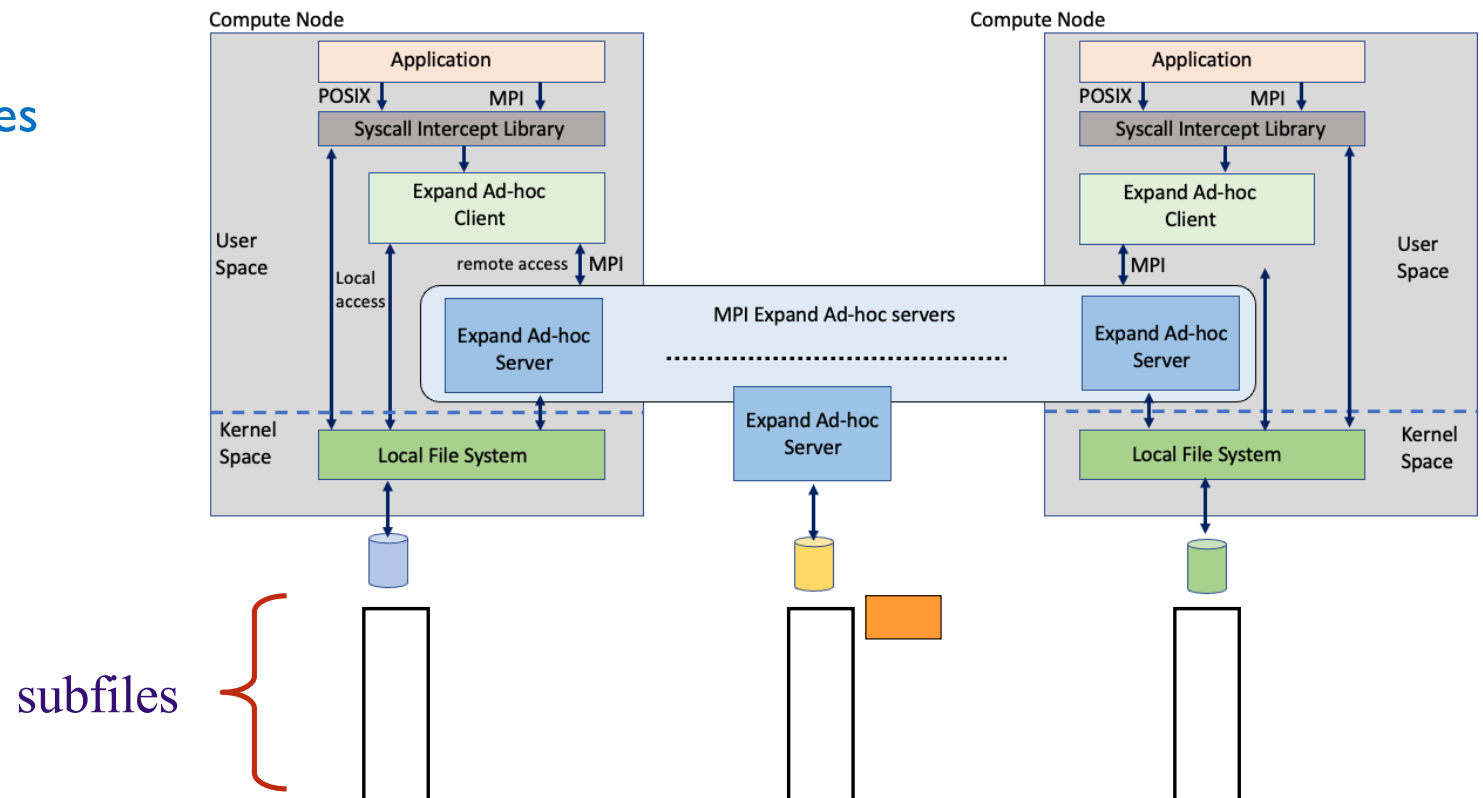
File in Expand



Data distribution

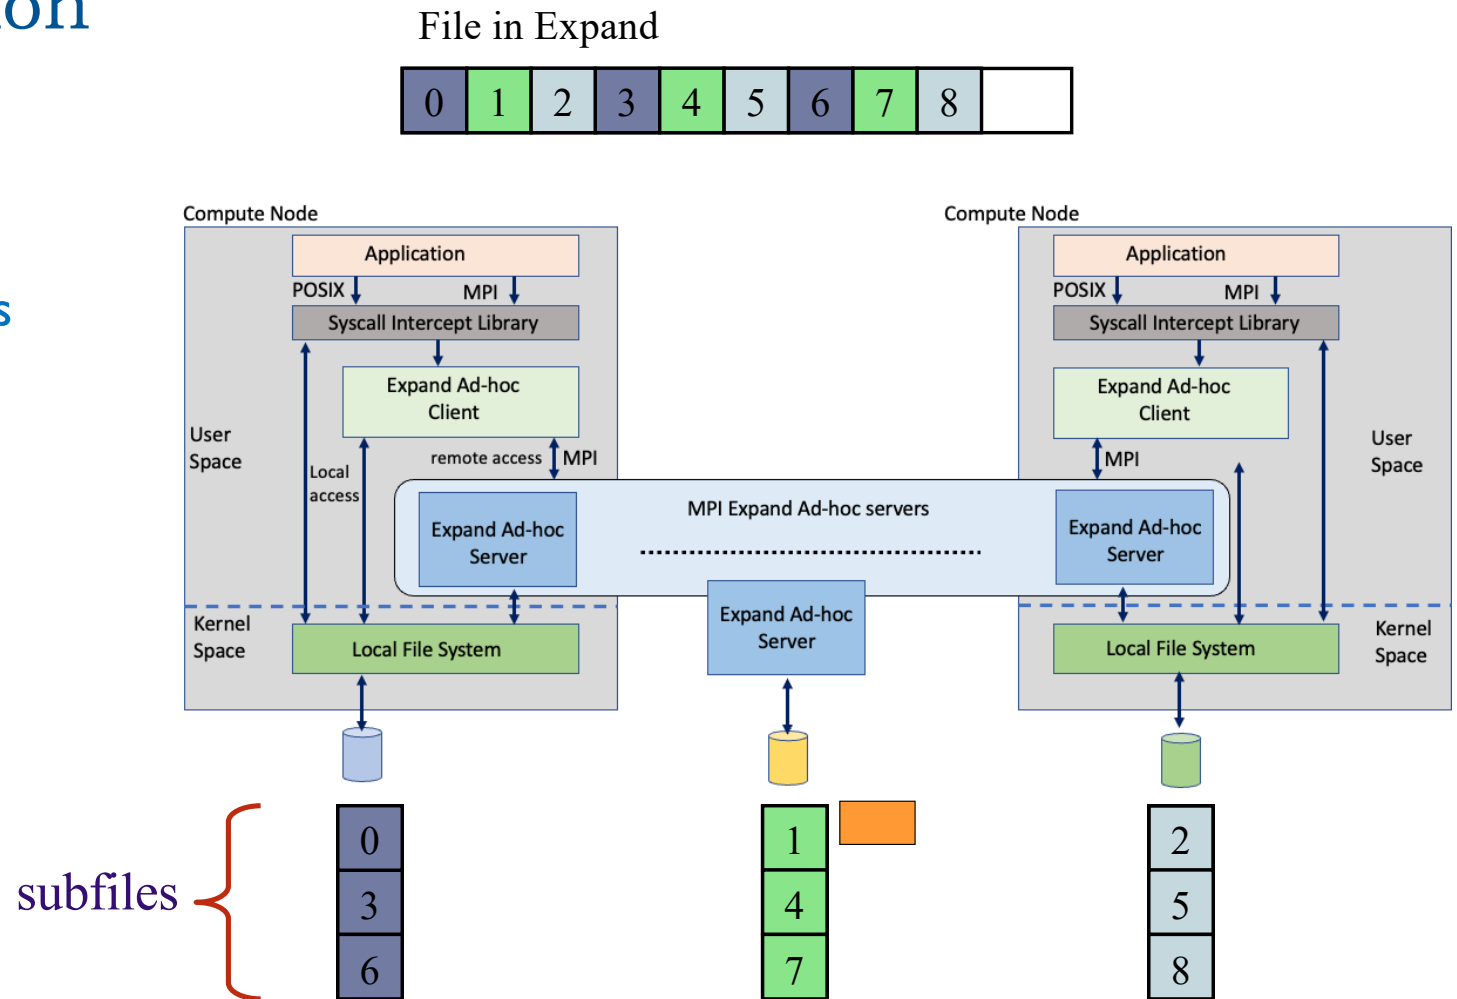
- ▶ File in Expand:
 - ▶ Metadata subfile
 - ▶ Several data subfiles
- ▶ Data distributed on different servers

File in Expand

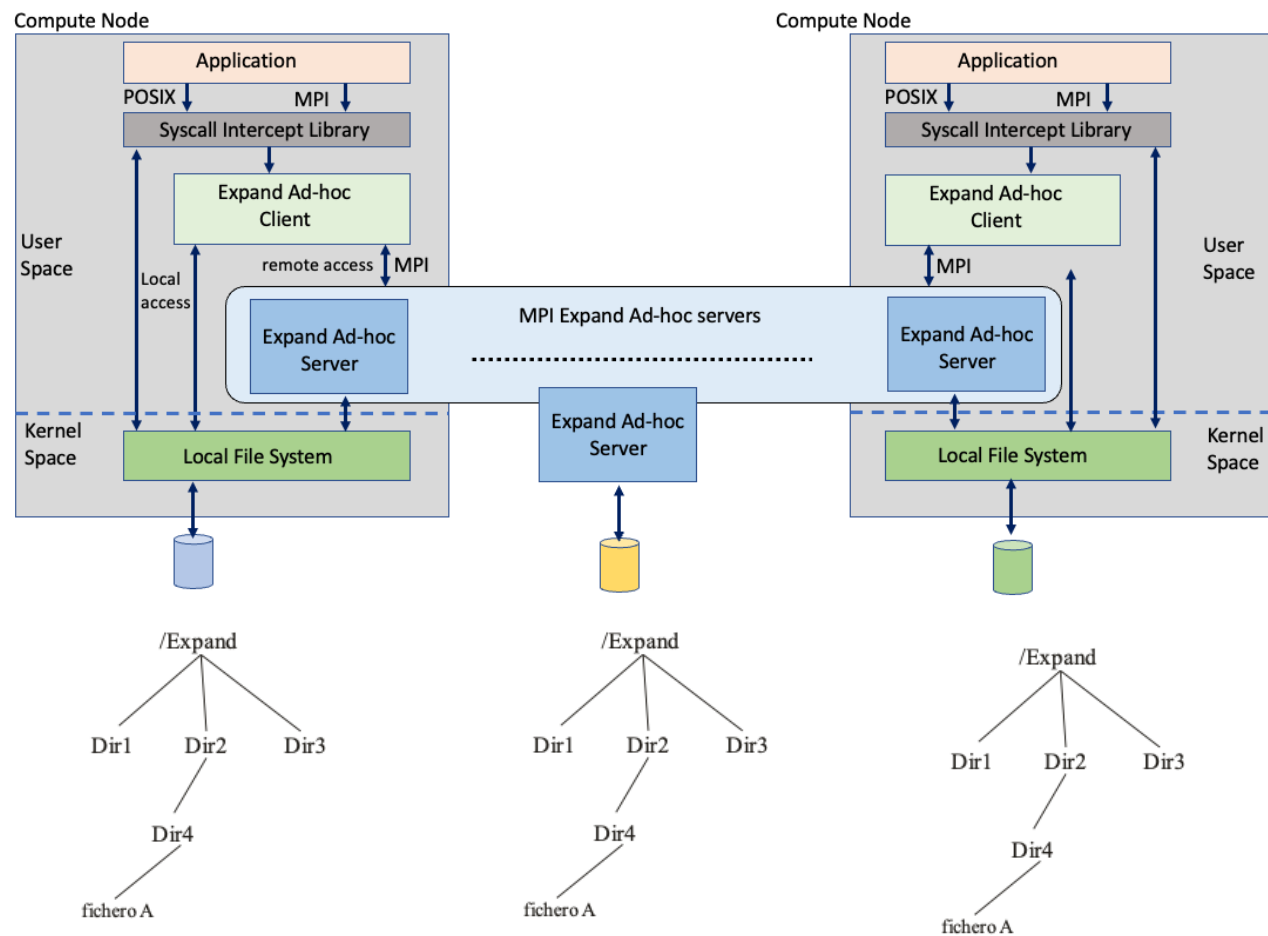


Data distribution

- ▶ File in Expand:
 - ▶ Metadata subfile
 - ▶ Several data subfiles
- ▶ Data distributed on different servers



Directory structure



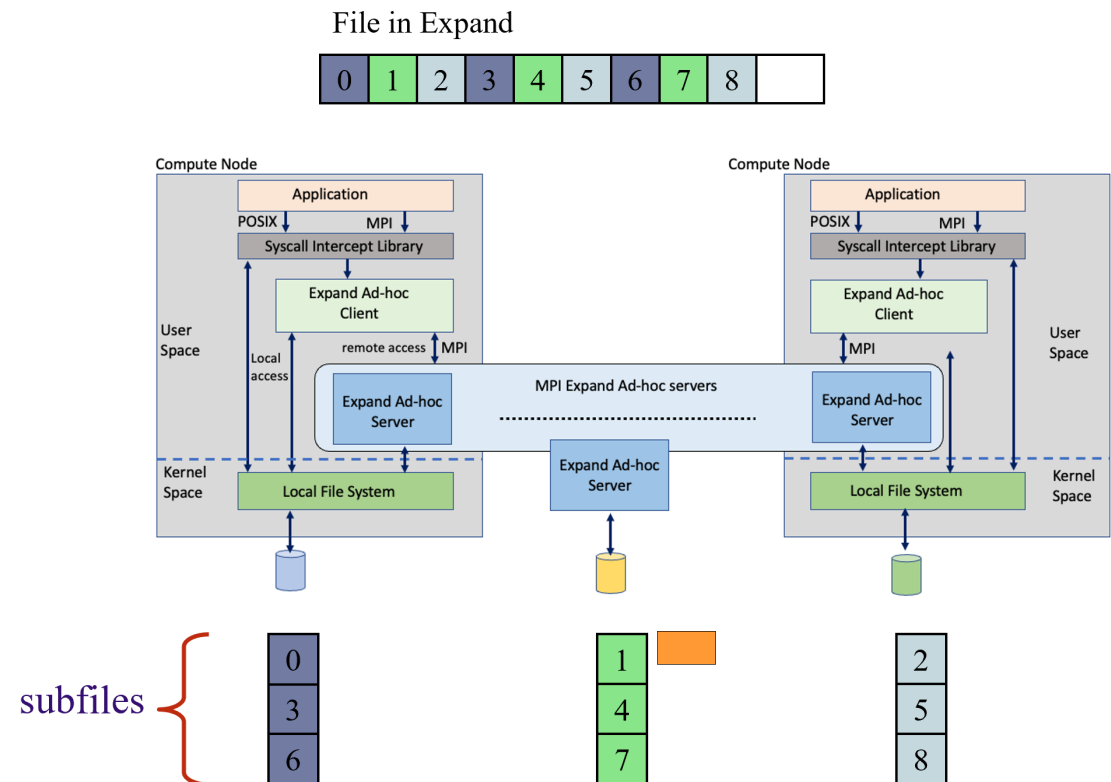
Metadata management

- ▶ Distributed metadata management:

- ▶ Two levels
- ▶ No locks
- ▶ No metadata manager

- ▶ Metadata distributed among servers:

- ▶ Master node
- ▶ Hash function(name)
- ▶ Load balancing



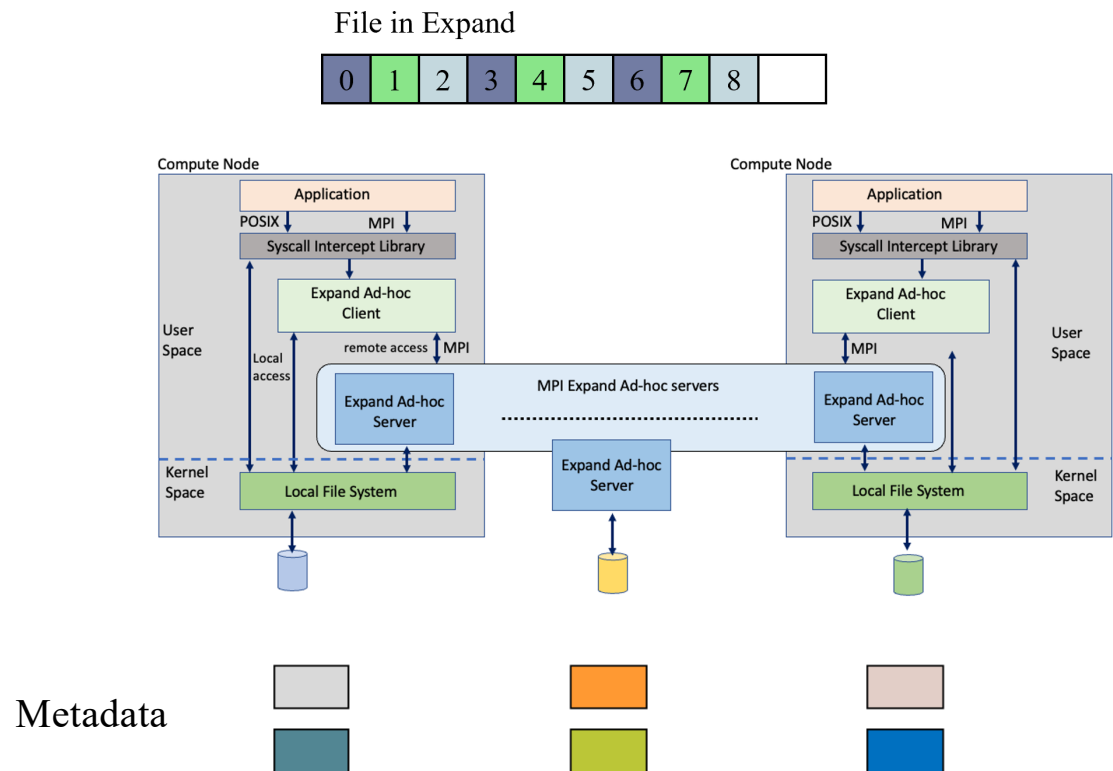
Metadata management

- ▶ Distributed metadata management:

- ▶ Two levels
- ▶ No locks
- ▶ No metadata manager

- ▶ Metadata distributed among servers:

- ▶ Master node
- ▶ Hash function(name)
- ▶ Load balancing



Evaluation

- ▶ **Benchmark used: IOR**
 - ▶ Open-source benchmark
 - ▶ A popular and effective way to evaluate the performance of distributed and parallel file systems is by using a variety of input/output (I/O) loads
- ▶ IOR access pattern: file sharing
 - ▶ All processes share the same file
 - ▶ Challenge for most parallel file systems

Evaluation

- ▶ Platform:
 - ▶ MareNostrum 4. Barcelona Supercomputing Center (BSC)
- ▶ Main properties:
 - ▶ *Nodes*: 3,456
 - ▶ *Total cores*: 165,888
 - ▶ *Main memory*: 384.75 TB
 - ▶ *SSD*: 240 GB
 - ▶ *Interconnection networks*: 100Gb Intel Omni-Path (Full-Fat Tree)
 - ▶ *Peak Performance*: 11.15 Petaflops
 - ▶ *Parallel file system*: GPFS

Evaluation

- ▶ File systems evaluated:

- ▶ **Expand Ad-Hoc** (512 KB of block size)
- ▶ **GekkoFS** (512 KB of block size)

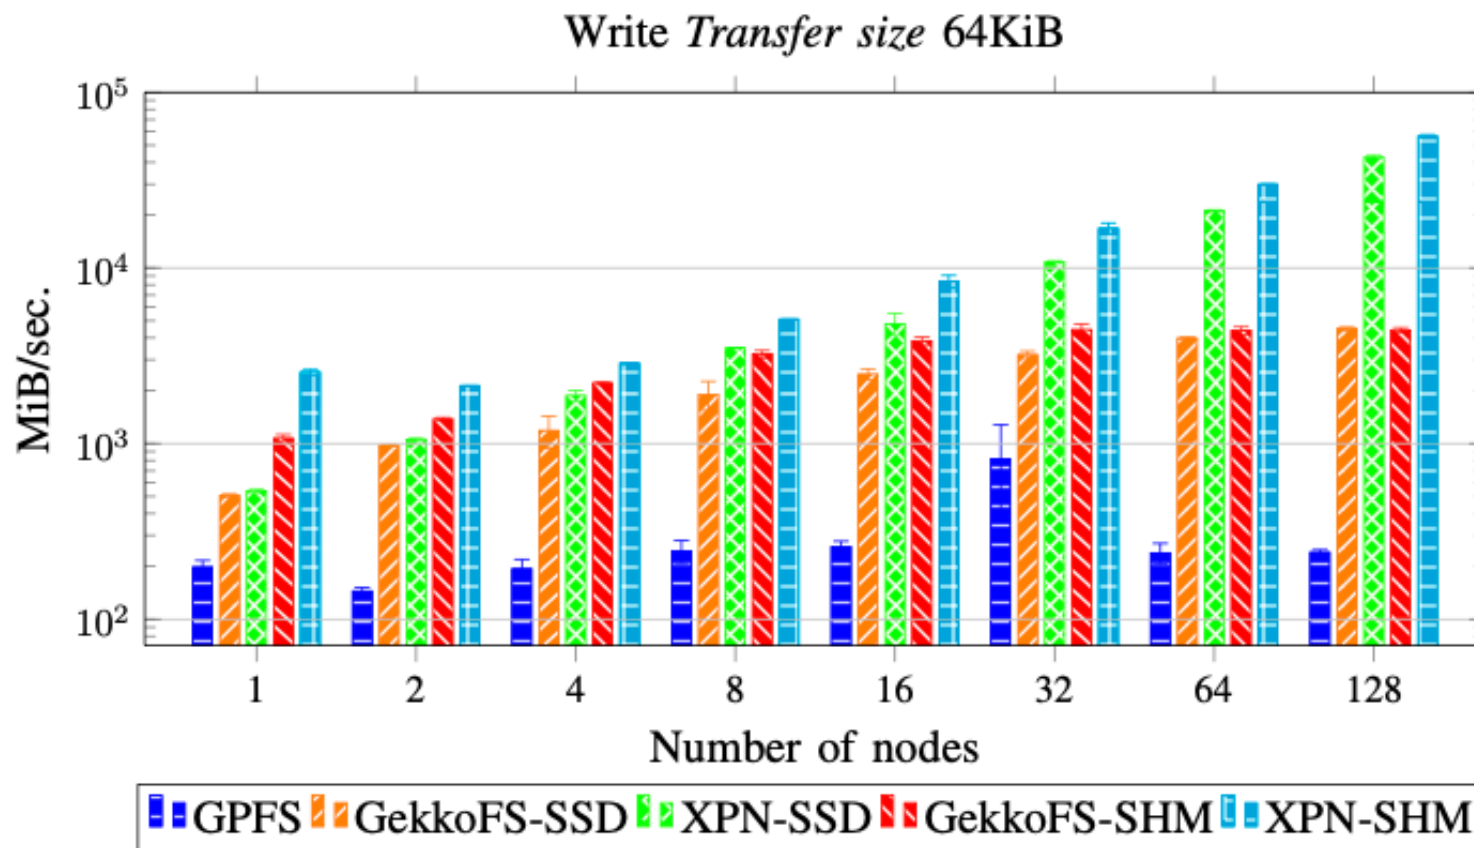
M.-A. Vef, N. Moti, T. Süß, T. Tocci, R. Nou, A. Miranda, T. Cortes, and A. Brinkmann, “GekkoFS: a temporary distributed file system for HPC applications,” in *2018 IEEE International Conference on Cluster Computing (CLUSTER)*. IEEE, 2018, pp. 319–324.

- ▶ **GPFS**: parallel file system used in MareNostrum 4

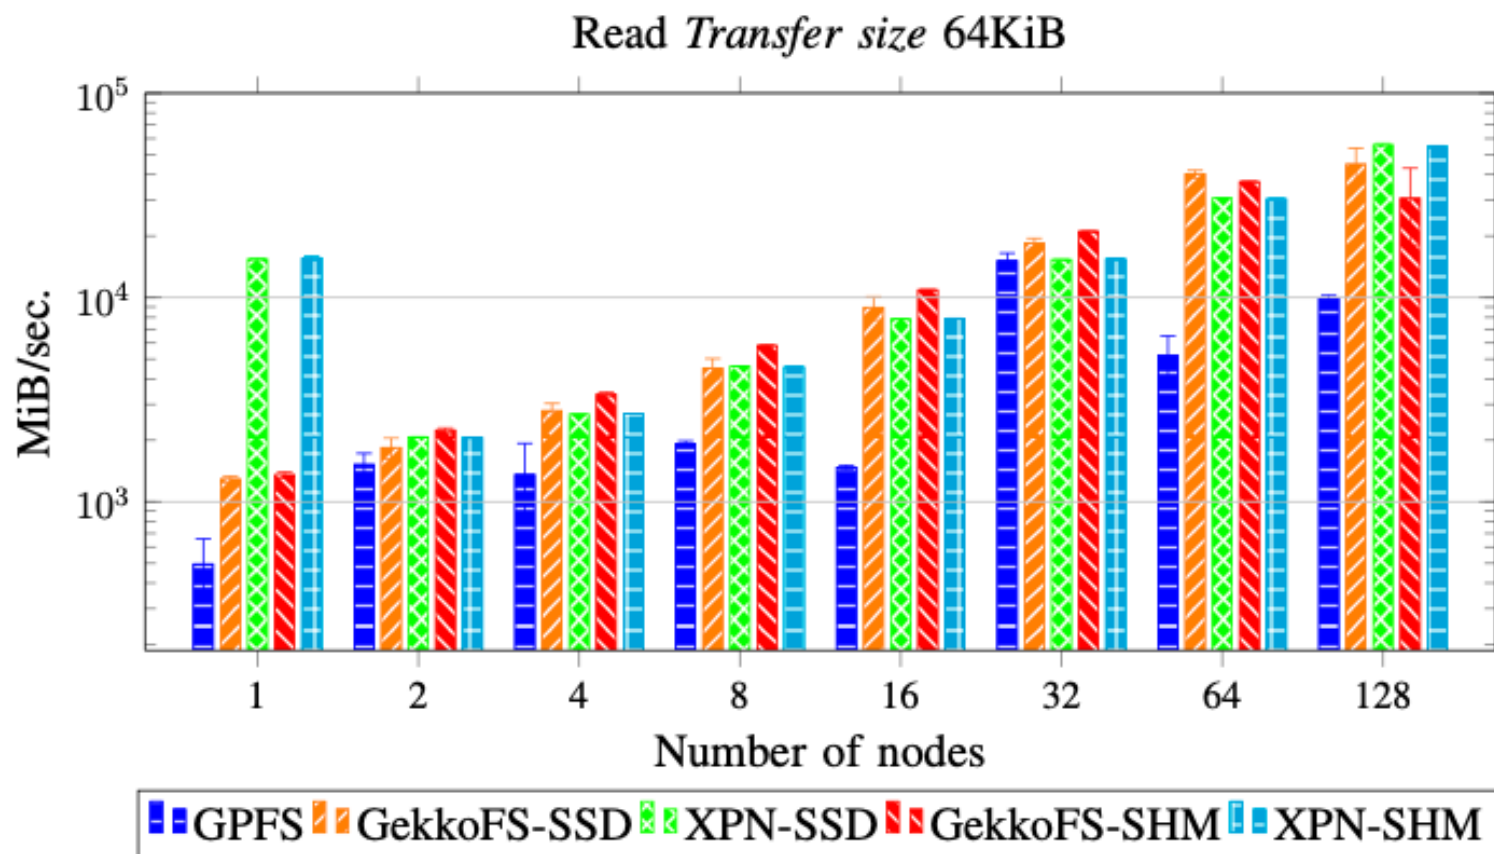
Evaluation

- ▶ Configuration:
 - ▶ Compute nodes: 1, 2, 4, ... up to 128 compute nodes
 - ▶ Local storage: SSD and Shared Memory (SHM) for GekkoFS and Expand
 - ▶ Transfer size used in IOR: 64 KiB, 512 KiB, and 1 MiB.
 - ▶ Client processes per compute node: 8
 - ▶ Operations: read and write in parallel on a shared file
 - ▶ Size written by each client: 4 GiB (resulting in a 4 TiB)
 - ▶ Number of processes: 8 up to 1024 processes
- ▶ All results in logarithmic scale
- ▶ All tests have been executed 10 times and the average results and standard deviation are shown in all figures.

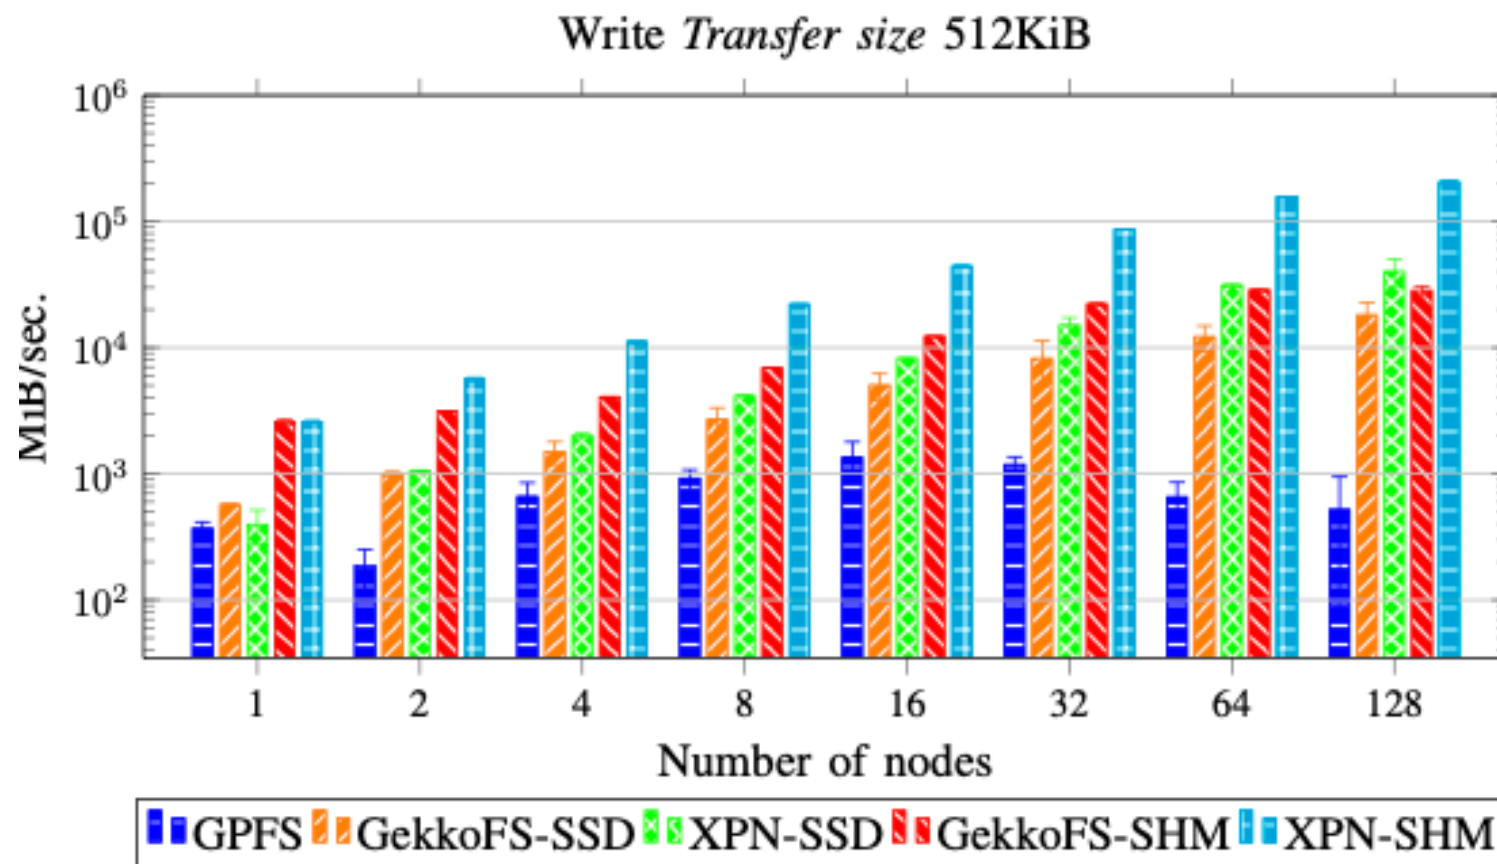
Results




Results



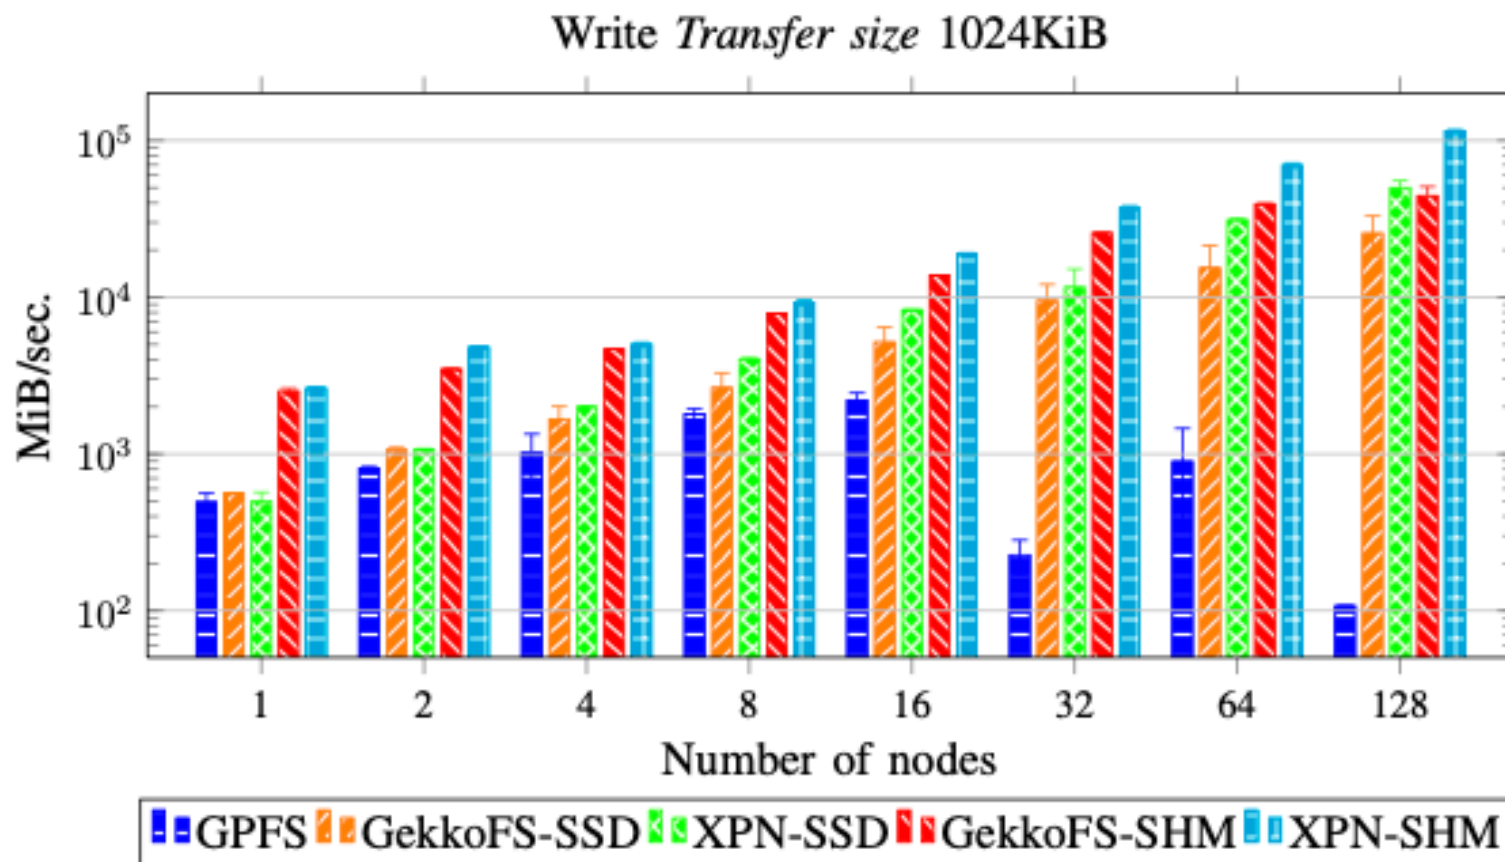
Results



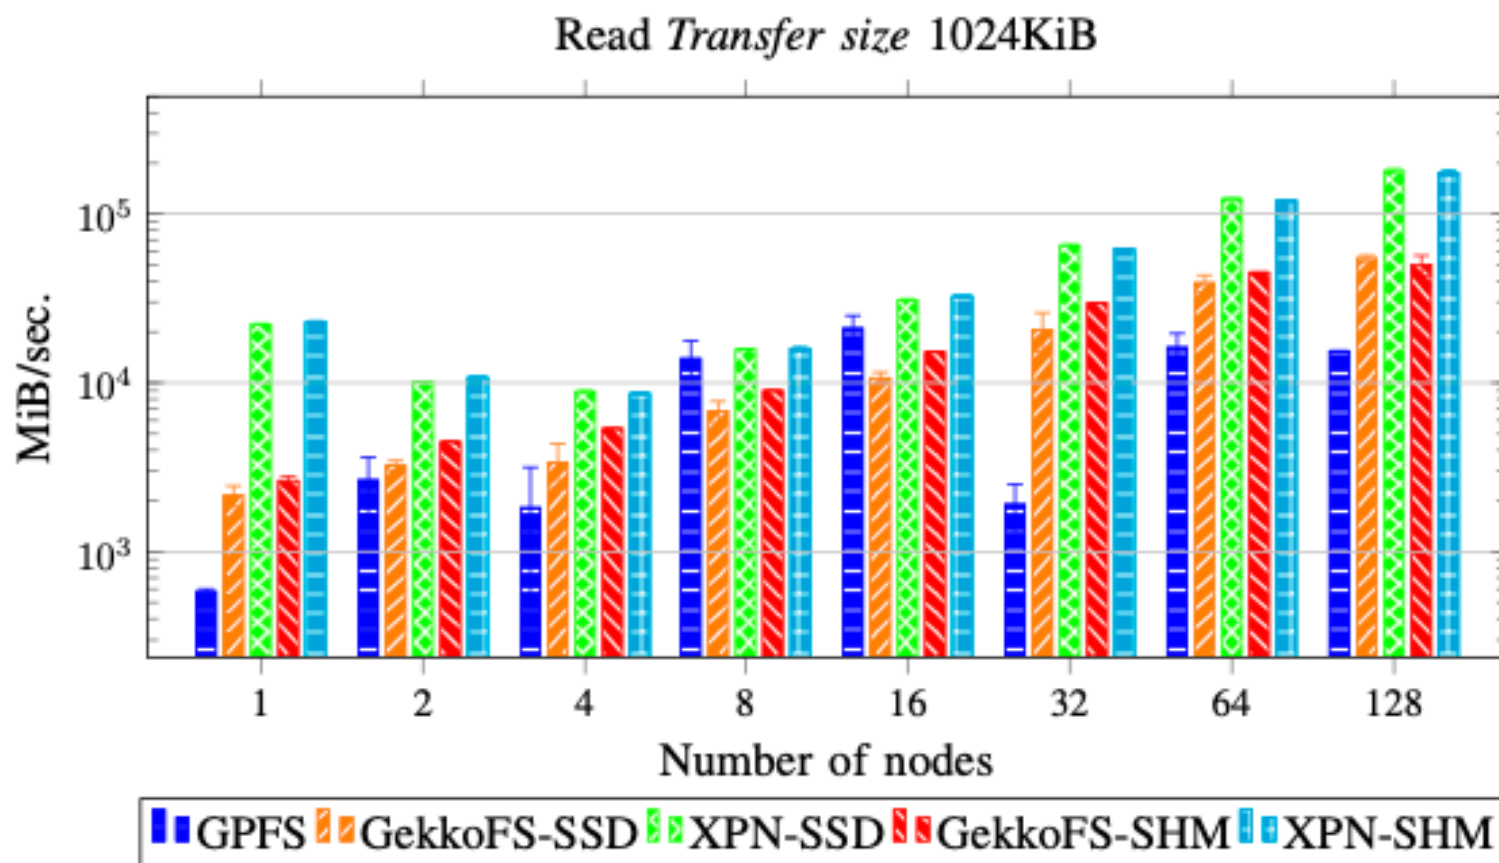
Results



Results



Results



Conclusions and Future Works

- ▶ Main conclusions:
 - ▶ Expand Ad-Hoc is available as an open-source project in the following GitHub repository:
 - ▶ <https://github.com/xpn-arcos/xpn>
 - ▶ Expand Ad-Hoc uses user-space MPI-based data servers on the same compute nodes on which applications are running
 - ▶ Expand Ad-Hoc provides a good performance compared to other solutions
- ▶ Future work:
 - ▶ Fault tolerant support

22nd IEEE International Symposium on Parallel and Distributed Computing

A new Ad-Hoc parallel file system for HPC environments based on the Expand parallel file system

Félix García Carballeira

felix.garcia@uc3m.es

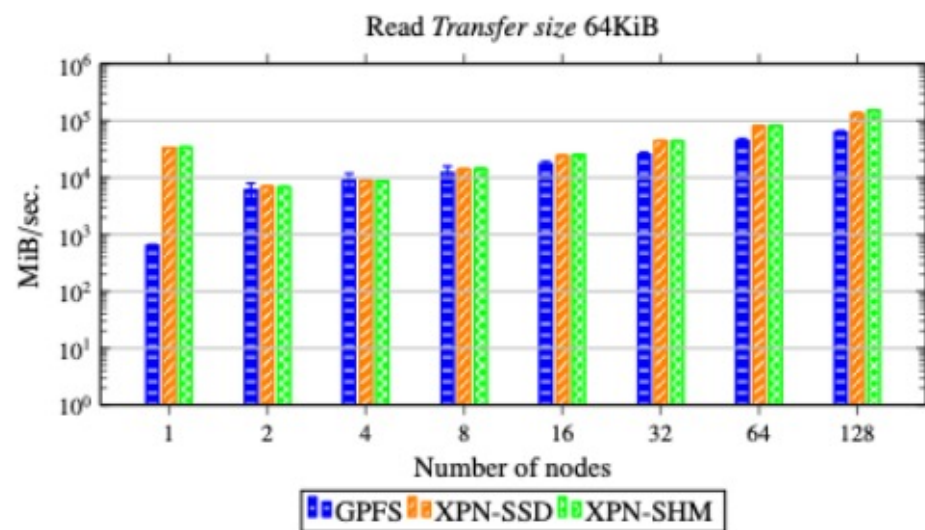
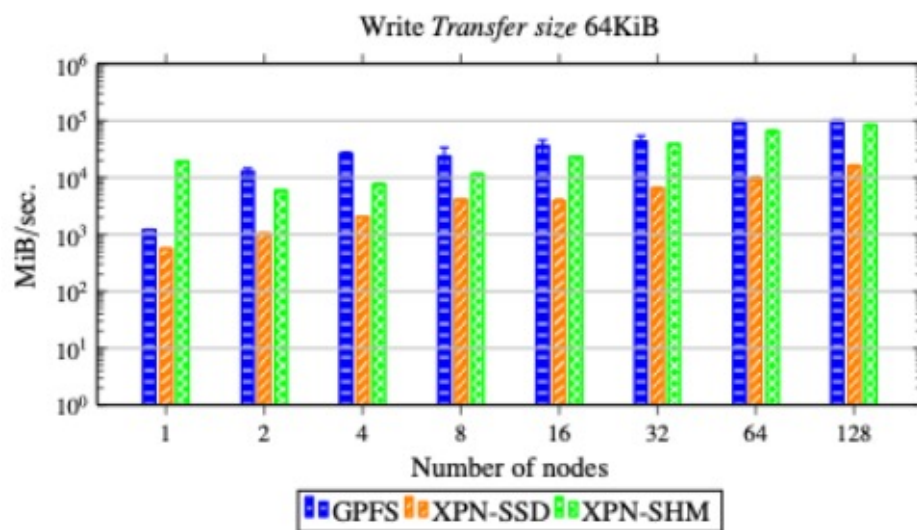


uc3m | Universidad **Carlos III** de Madrid

Evaluation

- ▶ Other results not included in the article:
 - ▶ A file per process. Each process accesses an individual file

Results



Results

